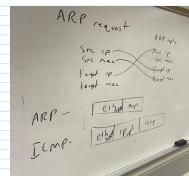
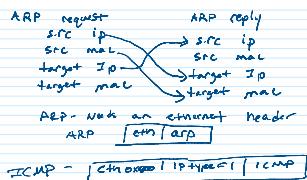


Consolidated Notes

Friday, October 28, 2022 6:05 PM



Internet Control Message Protocol (ICMP)

ICMP will have the following layout:

0 Eth type = 0x800 IP protocol = 1 ICMP type #

The ARP header (bytes 12-??) must follow the Ethernet (bytes 0-11). We must parse all the data to see what we have. We will take bytes 0-11 and copy them into the Eth header struct.

ICMP type #:

0 - echo reply (used to ping)

8 - echo request (used to ping)

Then take bytes 12 to 42 (for example) and copy that into an ARP struct.

ARP request have to broadcast everywhere

Eth	ARP
Source [source mac]	Operation (request or reply) set based on our needs
Destination [ffff:ff:ff:ff:ff:ff] (Broadcast mac address)	Fixed values for every ARP request/reply
Type = 0x806 (ARP Packet)	Hardware type = 1 (Ethernet)
	Hardware length = 6 bytes
	Protocol Type = 0x800
	Protocol Length = 4 bytes (for an IP address)
	Source IP Address [source IP]
	Source Hardware Address [source mac]
	Target Protocol Address [target IP]
	Target Hardware Address [target mac]

0 - 11 bytes

12 - ?? bytes

ARP Reply has to be unicast to who made the request

Eth	ARP
Source [target mac]	Operation (request or reply) set based on our needs
Destination [Source mac] (unicast)	Fixed values for every ARP request/reply
Type = 0x806 (ARP Packet)	Hardware type = 1 (Ethernet)
	Hardware length = 6 bytes
	Protocol Type = 0x800
	Protocol Length = 4 bytes (for an IP address)
	Source Protocol Address [target IP]
	Source Hardware Address [target mac]
	Target Protocol Address [source IP]
	Target Hardware Address [source mac]

0 - 11 bytes

12 - ?? bytes

Ignore all other packet types. No other packets should occur

Eth	ARP
Source mac	Operation (request or reply) set based on our needs
Dest mac	Fixed values for every ARP request/reply
Type = 0x806 (ARP Packet)	Hardware type = 1 (Ethernet)
	Hardware length = 6 bytes
	Protocol Type = 0x800
	Protocol Length = 4 bytes (for an IP address)
	Source IP Address [target IP]
	Source MAC Address [target mac]
	Target IP Address [source IP]
	Target MAC Address [source mac]

0 - 11 bytes

12 - ?? bytes

In this event we must send an ICMP "Host unreachable error message"

If the requested address does not exist we
the requester will not receive an ARP response.
The owner of the address is the only one that can
respond.

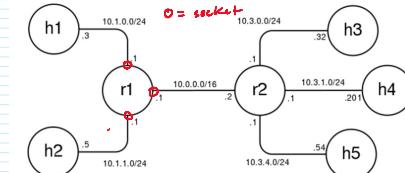
Request			
Eth	IP	ICMP Type	
Type = 0x800 (IP Packet)	IP Protocol = 1	Type = 8	
source mac	IP Protocol = 1	Checksum	
dest mac	if not 1 avoid it	id	
ip source	sequence	64 bits of data	
ip dest	id		
source mac	sequence		
dest mac	64 bits of data		
Protocol one means			
that ICMP is next			

A ping is an ICMP echo request packet

Note for obtaining mac address
iterate through if config
get addresses stuff

Not sure if we need this.

Ethernet header	IP header	ICMP header	User data	Ethernet CRC
https://org.rblab.ac.uk/ics301/queries/testpage/cmp.html				



Criteria	Points
Router accepts packets on packet sockets	5
Builds correct ARP response (including ethernet header)	10
Builds correct ICMP echo reply (including ethernet and ip headers)	10
Correctly uses packet socket to send responses as appropriate	5
All of the above work on all router interfaces	5

ARP http://www.arpipguide.com/free/ARPMessagedFormat.htm	
0	Hardware Type
4	Protocol Type
8	Length
12	Protocol Length
16	Opcode
20	
24	
28	
32	

Field Name	Size (bytes)	Description																				
		Hardware Type: This field specifies the type of hardware used for the local network transmitting the ARP message; thus, it also identifies the type of addressing used. Some of the most common values for this field:																				
HRD	2	<table border="1"> <thead> <tr> <th>HRD Value</th> <th>Hardware Type</th> </tr> </thead> <tbody> <tr> <td>1</td> <td>Ethernet</td> </tr> <tr> <td>5</td> <td>IEEE 802.11 wireless</td> </tr> <tr> <td>7</td> <td>ARCNET</td> </tr> <tr> <td>15</td> <td>Frame Relay</td> </tr> <tr> <td>16</td> <td>Asynchronous Transfer Mode (ATM)</td> </tr> <tr> <td>17</td> <td>Point-to-point</td> </tr> <tr> <td>18</td> <td>Fibre Channel</td> </tr> <tr> <td>19</td> <td>Asynchronous Transfer Mode (ATM)</td> </tr> <tr> <td>20</td> <td>Serial Line</td> </tr> </tbody> </table>	HRD Value	Hardware Type	1	Ethernet	5	IEEE 802.11 wireless	7	ARCNET	15	Frame Relay	16	Asynchronous Transfer Mode (ATM)	17	Point-to-point	18	Fibre Channel	19	Asynchronous Transfer Mode (ATM)	20	Serial Line
HRD Value	Hardware Type																					
1	Ethernet																					
5	IEEE 802.11 wireless																					
7	ARCNET																					
15	Frame Relay																					
16	Asynchronous Transfer Mode (ATM)																					
17	Point-to-point																					
18	Fibre Channel																					
19	Asynchronous Transfer Mode (ATM)																					
20	Serial Line																					
PRO	2	Protocol Type: This field is the complement of the Hardware Type field, specifying the type of layer three addresses used in the message. For IPv4 addresses, this value is 2048 (0800 hex), which corresponds to the EtherType code for the Internet Protocol.																				
HLN	1	Hardware Address Length: Specifies how long hardware addresses are in this message. For Ethernet or other networks using IEEE 802 MAC addresses, the value is 6.																				
PLN	1	Protocol Address Length: Again, the complement of the preceding field, specifies how long protocol (layer three) addresses are in this message. For IPv4 addresses this value is of course 4.																				
OP	2	Opcode: This field specifies the nature of the ARP message being sent. The two values (1 and 2) are used for regular ARP messages; other values are also defined to support other protocols that use the ARP frame format, such as RARP, RQRP, and others.																				
SHA	Variable, equals value in HLN field	Sender Hardware Address: The hardware (layer two) address of the device sending this message (which is the IP datagram source device on a request, and the IP datagram destination on a reply).																				
SPA	Variable, equals value in HLN field	Sender Protocol Address: The IP address of the device sending this message.																				
THA	Variable, equals value in HLN field	Target Hardware Address: The hardware (layer two) address of the device this message is being sent to. This is the IP datagram destination device on a request, and the IP datagram source on a reply.																				
TPA	Variable, equals value in PLN field	Target Protocol Address: The IP address of the device this message is being sent to.																				

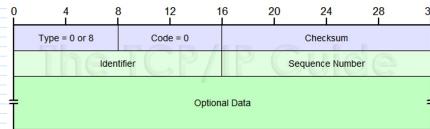


Table 99: ICMPv4 Echo and Echo Reply Message Format

Field Name	Size (bytes)	Description
Type	1	Type: Identifies the ICMP message type. For Echo messages the value is 8; for Echo Reply messages the value is 0.
Code	1	Code: Not used for Echo and Echo Reply messages; set to 0.
Checksum	2	Checksum: 16-bit checksum field for the ICMP header, as described in the topic on the ICMP common message format.
Identifier	2	Identifier: An identification field that can be used to help in matching Echo and Echo Reply messages.
Sequence Number	2	Sequence Number: A sequence number to help in matching Echo and Echo Reply messages.
Optional Data	Variable	Optional Data: Additional data to be sent along with the message (not specified).

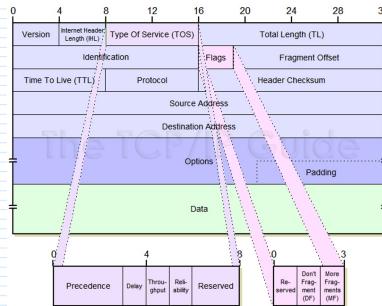


Table 66: Internet Protocol Version 4 (IPv4) Datagram Format

Field Name	Size (bytes)	Description
Version	4/2 (4 bits)	Version: Identifies the version of IP used to generate the datagram. For IPv4, this is of course the number 4. The purpose of this field is to ensure compatibility between devices that may be running different versions of IP. In general, a device running an older version of IP will reject datagrams created by newer implementations, under the assumption that the older version may not be able to interpret the newer datagram correctly.
IHL	4/2 (4 bits)	Internet Header Length (IHL): Specifies the length of the IP header in 32-bit words. This includes the length of any options fields and padding. The normal value of this field when no options are used is 5 (5 32-bit words). It can have a maximum value of 20 bytes. Contrast to the longer Total Length field below.
TOS	1	Type Of Service (TOS): A field designed to carry information to provide quality of service features, such as prioritized delivery, for IP datagrams. It was never widely used as originally defined, and its meaning has been subsequently redefined for use by a technique called Differentiated Services (DS). See below for more information.
TL	2	Total Length (TL): Specifies the total length of the IP datagram in bytes. Since this field is 16 bits wide, the maximum length of an IP datagram is 65,535 bytes, though most are much smaller.
Identification	2	Identification: This field contains a 16-bit value that is common to each of the fragments belonging to a particular message. For datagrams originally sent unfragmented it is still filled in, so it can be used if the datagram must be fragmented by a router during delivery. This field is used by the recipient to reassemble messages without accidentally mixing fragments from different messages. This is needed because fragments may arrive from multiple messages mixed together, since IP datagrams can be received out of order from any device. See the discussion of IP message fragmentation.

Flags	Flags: Three control flags, two of which are used to manage fragmentation (as described in the topic on fragmentation), and one that is reserved.		
	Sfield Name	Size (bytes)	Description
	Reserved	3/1(1/16)	Reserved: Not used.
DF	Don't Fragment	1/1(1/8)	Don't Fragment: When set to 1, specifies that the fragment must not be forwarded if it cannot be delivered in its current state. If a fragment is received in this state, the fragmentation process is generally "visible" to the recipient, who must then reassemble the message and start this flag again, however, used for forcing maximum transmission unit (MTU) of a link.
	MF	1/1(1/8)	More Fragments: When set to 1, indicates that there are more fragments in the message; when set to 0, indicates that this is the last fragment in the message. This flag is used in conjunction with the DF flag. If this flag is 1 (if fragmentation is used, and the DF flag is also set), the recipient knows when all fragments have been sent.
Fragment Offset	1/5/8 (13 bits)	Fragment Offset: When fragmentation occurs, this field specifies the offset, or position, in the overall message where the data in this fragment goes. It is specified in units of 8 bytes (64 bits). The first fragment has an offset of 0. Again, see the discussion of fragmentation for a description of how the field is used.	
TTL	1	Time To Live (TTL): Short version: Specifies how long the datagram is allowed to "live" on the network, in terms of router hops. Each router decrements the value of the TTL field (reduces it by one) prior to transmitting it. If the TTL field drops to zero, the datagram is assumed to have taken too long a route and is discarded. See below for the longer explanation of TTL.	

Protocol	Protocol: Identifies the higher-layer protocol (generally either a transport layer protocol or encapsulated version) being transported in the datagram. The values of this field were originally defined by the IETF (including TCP, UDP, and ICMP) and are now maintained by the Internet Assigned Numbers Authority (IANA).		
	Value (Hexadecimal)	Value (Decimal)	Protocol
00	0	Reserved	
01	1	ICMP	
02	2	IGMP	
03	3	GGP	
04	4	TCP-in-IP Encapsulation	
05	5	TCP	
06	6	ECN	
07	7	ESP	
11	17	UDP	
32	50	Encapsulating Security Payload (ESP) Extension Header	
33	51	Authentication Header (AH) Extension Header	

Note that the last two entries are used when IPsec inserts additional headers into the datagram; the AH or ESP headers.

Header Checksum: A checksum computed over the header to provide basic protection against corruption in transmission. This is not the more complex CRC code typically used by data link layer technologies such as Ethernet; it's just a 16-bit checksum. It is calculated by dividing the header bytes into words (a word is two bytes) and then adding them together. The data is not checksummed, only the header. At each hop the device receiving the datagram does the same checksum calculation and on a mismatch, discards the datagram as damaged.

Source Address: The 32-bit IP address of the originator of the datagram. Note that even though intermediate devices such as routers may handle the datagram, they do not normally put their address into this field—it is always the device that originally sent the datagram.

Destination Address: The 32-bit IP address of the intended recipient of the datagram. Again, even though devices such as routers may be the intermediate targets of the datagram, this field is always for the ultimate destination.

Options: One or more of several types of options may be included after the standard headers in certain IP datagrams. I discuss them in the topic that follows this one.

Padding: If one or more options are included, and the number of bits used for them is not a multiple of 32, enough zero bits are added to "pad out" the header to a multiple of 32 bits (4 bytes).

Data: The data to be transmitted in the datagram, either an entire higher-layer message or a fragment of one.

64 - 1518 byte

Ethernet Header (14 byte)



IEEE 802.3 Ethernet Frame Format

<https://markxover.github.io/blog/2020/05/09/networking/>

Library Document

Ethernet

<https://sites.uclouvain.be/SystInfo/usr/include/net/ether.h.html>

```
/* 10Mb/s ethernet header */
struct ether_header
{
    u_int8_t ether_dhost[ETH_ALEN]; /* destination eth addr */
    u_int8_t ether_shost[ETH_ALEN]; /* source ether addr */
    u_int16_t ether_type; /* packet type ID field */
} __attribute__((__packed__));
```

From <<https://sites.uclouvain.be/SystInfo/usr/include/net/ether.h.html>>

ARP + Ethernet

https://sites.uclouvain.be/SystInfo/usr/include/netinet/if_ether.h.html

```
struct ether_arp {
    struct arphdr ea_hdr; /* fixed-size header */
    u_int8_t arp_sha[ETH_ALEN]; /* sender hardware address */
    u_int8_t arp_spa[4]; /* sender protocol address */
    u_int8_t arp_tha[ETH_ALEN]; /* target hardware address */
    u_int8_t arp_tpa[4]; /* target protocol address */
};
```

From <https://sites.uclouvain.be/SystInfo/usr/include/netinet/if_ether.h.html>

ARP struct that's used in the ARP+Eth header

https://sites.uclouvain.be/SystInfo/usr/include/net/if_arp.h.html

```
struct arphdr
{
    unsigned short int ar_hrd; /* Format of hardware address */
    unsigned short int ar_pro; /* Format of protocol address */
    unsigned char ar_hln; /* Length of hardware address */
    unsigned char ar_pln; /* Length of protocol address */
    unsigned short int ar_op; /* ARP opcode (command) */
};
```

We don't need to use this at all.

```
#if 0
/* Ethernet looks like this : This bit is variable sized
however... */
unsigned char __ar_sha[ETH_ALEN]; /* Sender hardware address */
unsigned char __ar_sip[4]; /* Sender IP address */
unsigned char __ar_tha[ETH_ALEN]; /* Target hardware address */
unsigned char __ar_tip[4]; /* Target IP address */
#endif
};
```

From <https://sites.uclouvain.be/SystInfo/usr/include/net/if_arp.h.html>

ICMP

https://sites.uclouvain.be/SystInfo/usr/include/netinet/ip_icmp.h.html

```
struct icmpHdr
{
    u_int8_t type; /* message type */
    u_int8_t code; /* type sub-code */
    u_int16_t checksum;
    union
    {
        struct
        {
            u_int16_t id; /* We don't need to use this. This is already factored
                           in while using this struct */
            u_int16_t sequence;
        } echo; /* echo datagram */
        u_int32_t gateway; /* gateway address */
        struct
        {
            u_int16_t __unused;
            u_int16_t mtu;
        } frag; /* path mtu discovery */
        u_int8_t un;
    };
};
```

From <https://sites.uclouvain.be/SystInfo/usr/include/netinet/ip_icmp.h.html>

IP

<https://sites.uclouvain.be/SystInfo/usr/include/netinet/ip.h.html>

```
struct iphldr /* We don't have to change this. The system knows what to do.
{
#if __BYTE_ORDER == __LITTLE_ENDIAN
    unsigned int ihl:4;
    unsigned int version:4;
#elif __BYTE_ORDER == __BIG_ENDIAN
    unsigned int version:4;
    unsigned int ihl:4;
#else
    # error "Please fix <bits/endian.h>"
```

The options start here.

```
#endif
    u_int8_t tos;
    u_int16_t tot_len;
    u_int16_t id;
    u_int16_t frag_off;
    u_int8_t ttl;
    u_int8_t protocol;
    u_int16_t check;
    u_int32_t saddr;
    u_int32_t daddr;
};
```

From <<https://sites.uclouvain.be/SystInfo/usr/include/netinet/ip.h.html>>